

Winter Semester 2025-2026

Name: Panagiotis Kalogiannis

ID: csd4888

Date: January 17, 2026

Cloud Resource Allocation Project Report

1. Introduction

This report presents the implementation and results of a game-theoretic resource allocation system for cloud computing services. The project implements a two-step optimization process where three users (tasks) compete for five computational resources provided by a cloud provider.

2. Configuration

2.1 Student ID and Weight Calculation

Student ID Last Two Digits (D): 88

Weight Calculations:

Task	Weight Time (wt)	Weight Expense (we)
S1	$D/100 = 0.88$	$1 - wt = 0.12$
S2	$(D+1)\%100/100 = 0.89$	$1 - wt = 0.11$
S3	$(D+2)\%100/100 = 0.90$	$1 - wt = 0.10$

2.2 Task Configuration

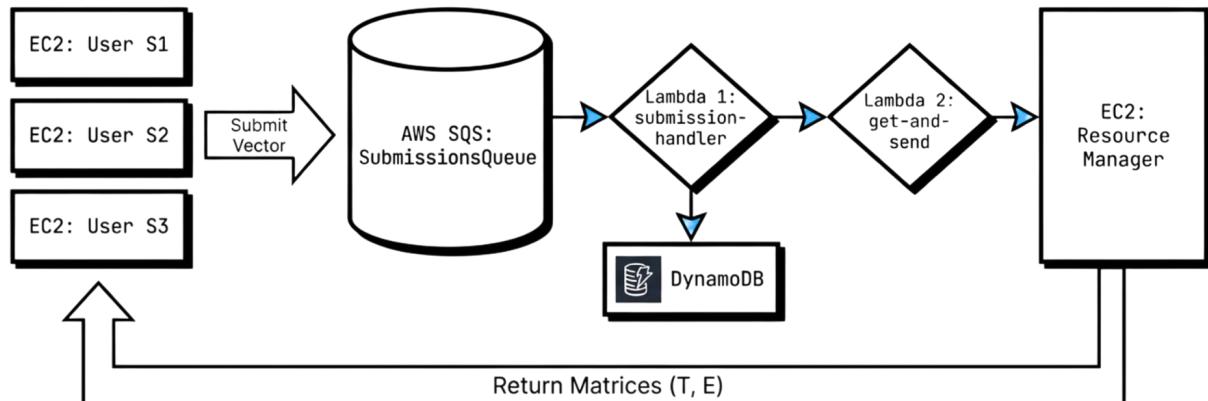
Task	Subtasks	Deadline (s)	Budget (€)
S1	2	500	20
S2	3	300	30
S3	4	800	30

2.3 Resource Configuration

Resource	Price (€/s)	Execution Times (S1, S2, S3)
R1	1.0	5.0, 6.0, 4.0
R2	1.2	4.2, 5.0, 3.5
R3	1.5	3.6, 4.0, 3.2
R4	1.8	3.0, 3.5, 2.8
R5	2.0	2.8, 3.0, 2.4

3. Architecture

3.1 System Architecture Diagram



Workflow: Asynchronous Event-Driven Architecture. Users optimize locally → Submit to Queue → Aggregated in DynamoDB → Batch Processed by Manager → Results Distributed.

Architecture Flow Description:

1. **Submission Phase:** Users (S1, S2, S3) perform local optimization and submit allocation vectors to SQS queue
2. **Processing Phase:** Lambda 1 (submission-handler) is triggered by SQS messages, stores allocations in DynamoDB

3. **Aggregation Phase:** When all 3 users have submitted, Lambda 1 triggers Lambda 2 (get-and-send)
4. **Calculation Phase:** Lambda 2 retrieves allocations from DynamoDB and sends them to Resource Manager
5. **Distribution Phase:** Resource Manager calculates actual execution times and expenses (accounting for multiplexing) and sends T and E matrices back to all users
6. **Status Update:** Lambda 2 updates allocation status in DynamoDB to 'processed'

The system consists of:

1. **User Layer:** 3 EC2 instances running Flask applications (one per user)
2. **Message Queue:** SQS queue for asynchronous communication
3. **Processing Layer:** 2 Lambda functions for workflow orchestration
4. **Storage Layer:** DynamoDB for persistent allocation storage
5. **Provider Layer:** Resource Manager EC2 instance for matrix calculations

3.2 Component Description

User EC2 Instances:

- Perform brute-force optimization to find best allocation
- Submit allocation vectors to SQS
- Receive and process T and E matrices from provider
- Calculate actual utilities

Lambda 1 (Submission Handler):

- Triggered by SQS messages
- Stores allocations in DynamoDB
- Checks if all users submitted
- Triggers Lambda 2 when ready

Lambda 2 (Get and Send):

- Retrieves allocations from DynamoDB
- Sends to Resource Manager
- Updates processing status

Resource Manager EC2:

- Receives allocation vectors from all users
- Calculates actual execution times (with multiplexing)
- Calculates expenses
- Returns T and E matrices to users

4. Implementation

4.1 AWS Resources Created

The screenshot shows the AWS EC2 Services page. On the left, there's a sidebar with links to Services, Features, Documentation, and Tutorials. The main area lists EC2 Instances, EC2 Resource Health, and a Dashboard. A bar chart titled 'Cost (US\$)' shows spending over time, with a significant peak in December 2025. A 'Create application' button is visible at the top right.

The screenshot shows the 'Launch an instance' page. It includes fields for 'Name and tags', 'Application and OS Images (Amazon Machine Image)', and 'Amazon Machine Image (AMI)'. The AMI selected is 'Ubuntu Server 24.04 LTS (HVM), SSD Volume Type'. The 'Launch instance' button is highlighted in orange at the bottom right.

EC2 Instances:

- User-S1: 18.207.117.128 (ec2-18-207-117-128.compute-1.amazonaws.com)

Instance summary for i-0a5d3fa0e1669b48c (My-FirstEC2-Instance)

Attribute	Value
Public IP4 address	18.207.117.128 [open address]
Instance state	Running
Private IP DNS name (IPv4 only)	ip-172-31-29-232.ec2.internal
Instance type	t2.micro
VPC ID	vpc-0d423e8abb901bba7
Subnet ID	subnet-03e45382800e5767
Instance ARN	arn:aws:ec2:us-east-1:210058569281:instance/i-0a5d3fa0e1669b48c
Elastic IP addresses	-
AWS Compute Optimizer finding	Opt-in to AWS Compute Optimizer for recommendations. [Learn more]
Auto Scaling Group name	-
Managed	Managed

Details Tab

Instance details

Attribute	Value
AMI ID	ami-0ecb62995f68bb549
AMI name	ubuntu/images/hvm-ssd-gp3/ubuntu-noble-24.04-amd64-server-20251022
Stop protection	Disabled
Launch time	Sat Jan 17 2026 17:49:37 GMT+0500 (Pakistan Standard Time) (2 minutes)

Platform details

- Platform: Linux/UNIX
- Termination protection: Disabled
- AMI location: <https://amazon/ubuntu/images/hvm-ssd-gp3/ubuntu-noble-24.04-amd64-server-20251022>

Terminal Session (Ubuntu 24.04.3 LTS)

```

$ cv project chmod 400 cs452-project-key.pem
$ cv project chmod 400 ~Downloads/cs452-project-key.pem
$ cv project ssh -i cs452-project-key.pem ubuntu@18.207.117.128
The authenticity of host '18.207.117.128 (18.207.117.128)' can't be established.
ED25519 key fingerprint is SHA256:05520D01Er/CGCIHwgep8XXV80AC80ghRFIP6zMIu8.
This host key is known by the following other names/addresses:
  -> /ssh/known_hosts:22: ec2-18-207-117-128.compute-1.amazonaws.com
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added '18.207.117.128' (ED25519) to the list of known hosts.
Welcome to Ubuntu 24.04.3 LTS (GNU/Linux 6.14.0-1015-aws x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/pro

System information as of Sat Jan 17 12:55:30 UTC 2026

System load: 0.08      Processes:          109
Usage of /: 9.4% of 18.33GB   Users logged in:     0
Memory usage: 22%           IPv4 address for enX0: 172.31.29.232
Swap usage:  0%           Swap space available: 0B

Expanded Security Maintenance for Applications is not enabled.

0 updates can be applied immediately.

Enable ESM Apps to receive additional future security updates.
See https://ubuntu.com/esm or run: sudo pro status

The list of available updates is more than a week old.
To check for new updates run: sudo apt update

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

ubuntu@ip-172-31-29-232:~$ 
```

Screenshot of the AWS EC2 Security Groups page showing the successful modification of a security group rule.

Details

Inbound security group rules successfully modified on security group (sg-0eda0e1b30cf69e3a | launch-wizard-3)

sg-0eda0e1b30cf69e3a - launch-wizard-3

Actions

Details

Security group name launch-wizard-3	Security group ID sg-0eda0e1b30cf69e3a	Description launch-wizard-3 created 2026-01-17T12:41:44.528Z	VPC ID vpc-0d423e8abb901bba7.t
Owner 210058569281	Inbound rules count 4 Permission entries	Outbound rules count 1 Permission entry	

Inbound Rules | Outbound rules | Sharing | VPC associations | Tags

Inbound rules (4)

Name	Security group rule ID	IP version	Type	Protocol	Port range	Source	Description
-	sgr-0dfec7e9b3d4b2f3f	IPv4	HTTP	TCP	80	0.0.0.0/0	-
-	sgr-0b965fd6953eac7da4	IPv4	HTTPS	TCP	443	0.0.0.0/0	-
-	sgr-0b2649d1142e456c4	IPv4	SSH	TCP	22	0.0.0.0/0	-
-	sgr-0e32cd8a7f59e8d	IPv4	Custom TCP	TCP	5000	0.0.0.0/0	Flask app

CloudShell | Feedback

Problems Output Debug Console Terminal Ports

```

numpy installed successfully
(venv) ubuntu@ip-172-31-29-232:~/app$ export USER_ID=S1
(venv) ubuntu@ip-172-31-29-232:~/app$ export PORT=5000
(venv) ubuntu@ip-172-31-29-232:~/app$ nohup python user/user_app.py > user.log 2>&1 &
[1] 3358
(venv) ubuntu@ip-172-31-29-232:~/app$ cat user.log
nohup: ignoring input

*** Initialized Optimizer for S1 ***
Subtasks: 2
Weights: wt=0.88, we=0.12
Constraints: T≤500s, M≤20€

=====
Starting User Flask App: S1
Port: 5000
=====

* Serving Flask app 'user_app'
* Debug mode: on
WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
* Running on all addresses (0.0.0.0)
* Running on http://127.0.0.1:5000
* Running on http://172.31.29.232:5000
Press CTRL+C to quit
* Restarting with stat
* Debugger is active!
* Debugger PIN: 231-396-925
(venv) ubuntu@ip-172-31-29-232:~/app$ ps aux | grep user_app.py
ubuntu  3358  0.4  5.5 166228 54744 pts/0   S  13:07  0:00 python user/user_app.py
ubuntu  3359  0.7  5.5 23972 54696 pts/0   Sl  13:07  0:00 /home/ubuntu/venv/bin/python user/user_app.py
ubuntu  3360  0.0  0.2  7084  2076 pts/0   S+ 13:08  0:00 grep --color=auto user_app.py
(venv) ubuntu@ip-172-31-29-232:~/app$ curl http://localhost:5000/health
{
    "status": "healthy",
    "step": 0,
    "user_id": "S1"
}
(venv) ubuntu@ip-172-31-29-232:~/app$ 
```

© 2026, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookies preferences

- User-S2: 50.19.166.35 (ec2-50-19-166-35.compute-1.amazonaws.com)

EC2 > Instances > i-0e00e90103a3d52c0

Instance summary for i-0e00e90103a3d52c0 (S2-User)

Updated 3 minutes ago	Public IP4 address 50.19.166.35 open address	Private IP4 addresses 172.31.17.102
Instance ID i-0e00e90103a3d52c0	Instance state Running	Public DNS ec2-50-19-166-35.compute-1.amazonaws.com open address
IPv6 address -	Private IP DNS name (IPv4 only) ip-172-31-17-102.ec2.internal	Elastic IP addresses -
Hostname type IP name: ip-172-31-17-102.ec2.internal	Instance type t2.micro	AWS Compute Optimizer finding Opt-in to AWS Compute Optimizer for recommendations. Learn more
Answer private resource DNS name IPv4 (A)	VPC ID vpc-0d423e8abb901bba7	Auto Scaling Group name -
Auto-assigned IP address 50.19.166.35 [Public IP]	Subnet ID subnet-034e45382800e5767	Managed false
IAM Role LabRole	Instance ARN arn:aws:ec2:us-east-1:21005869281:instance/i-0e00e90103a3d52c0	Platform details Linux/UNIX
IMDv2 Required	Launch time Sat, Jan 17, 2026 18:17:56 GMT+0500 (Pakistan Standard Time) (less than a minute ago)	Termination protection Disabled
Operator -	AMI location amazon/ubuntu/images/hvm-ssd/ubuntu-noble-24.04-amd64-s	AMI location amazon/ubuntu/images/hvm-ssd/ubuntu-noble-24.04-amd64-s

Details **Status and alarms** **Monitoring** **Security** **Networking** **Storage** **Tags**

Instance details

Inbound rules

Edit inbound rules

Inbound rules

Type	Protocol	Port range	Source	Description - optional
HTTP	TCP	80	Custom	
SSH	TCP	22	Custom	
HTTPS	TCP	443	Custom	
Custom TCP	TCP	5000	Anywhere	Flask Project Instance 2 S2 User

Add rule

Rules with source of 0.0.0.0 or ::/0 allow all IP addresses to access your instance. We recommend setting security group rules to allow access from known IP addresses only.

Cancel **Preview changes** **Save rules**

```

Successfully installed Flask-3.1.2 blinker-1.9.0 boto3-1.42.30 botocore-1.42.30 certifi-2026.1.4 charset_normalizer-3.4.4 click-8.3.1 idna-3.11 itsdangerous-2.2.0 jinja2-3.1.6 jmespath-1.0.1 markupsafe-3.0.3 numpy-2.4.1 python-dateutil-2.9.0.post0 requests-2.32.5 s3transfer-0.16.0 six-1.17.0 urllib3-2.6.3 werkzeug-3.1.5
(venv) ubuntu@ip-172-31-17-102:~$ mkdir -p ~/app
(venv) ubuntu@ip-172-31-17-102:~/app$ export USER_ID=S2
(venv) ubuntu@ip-172-31-17-102:~/app$ export PORT=5000
(venv) ubuntu@ip-172-31-17-102:~/app$ nohup python user/user_app.py > user.log 2>&1 &
[1] 3191
(venv) ubuntu@ip-172-31-17-102:~/app$ cat user.log
nohup: ignoring input
== Initialized Optimizer for S2 ==
Subtasks: 3
Weights: wt=0.89, we=0.11
Constraints: T=300s, M=300s
=====
Starting User Flask App: S2
Port: 5000
=====

* Serving Flask app 'user_app'
* Debug mode: on
WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
* Running on all addresses (0.0.0.0)
* Running on http://127.0.0.1:5000
* Running on http://172.31.17.102:5000
Press CTRL+C to quit
* Restarting with stat
* Debugger is active!
* Debugger PIN: 865-556-343
(venv) ubuntu@ip-172-31-17-102:~/app$ 

```

- User-S3: 98.80.12.101 (ec2-98-80-12-101.compute-1.amazonaws.com)

Instance summary for i-05e20216b65e8bca9 (S3-User)

Details | Status and alarms | Monitoring | Security | Networking | Storage | Tags

AMI ID	ami-0ecb62995f68bb549	Monitoring	Platform details
AMI name	ubuntu/images/hvm-ssd-gp3/ubuntu-noble-24.04-amd64-server-20251022	Allowed image	Linux/UNIX
Stop protection	Disabled	Launch time	Termination protection
		Sat Jan 17 2026 18:28:49 GMT+0500 (Pakistan Standard Time) (EST)	Disabled
AMI location			

```

Problems Output Debug Console Terminal Ports
● cv_project chmod 400 ~/Downloads/s3-user.pem
○ → cv_project ssh -i ~/Downloads/s3-user.pem ubuntu@98.80.12.101
The authenticity of host '98.80.12.101 (98.80.12.101)' can't be established.
ED25519 key fingerprint is SHA256:Q3dhH1nKcgw/Y+0ig+q+XgIMxEW7sqE7OTwyBb4KE.
This key is not known by any other host.
Are you sure you want to continue connecting (yes/no/fingerprint)? yes
Warning: Permanently added '98.80.12.101' (ED25519) to the list of known hosts.
Welcome to Ubuntu 24.04.3 LTS (GNU/Linux 6.14.0-1015-aws x86_64)

 * Documentation: https://help.ubuntu.com
 * Management: https://landscape.canonical.com
 * Support: https://ubuntu.com/pro

System information as of Sat Jan 17 13:32:21 UTC 2026
System load: 0.11 Processes: 114
Usage of /: 9.4% of 18.33GB Users logged in: 0
Memory usage: 22% IPv4 address for enx0: 172.31.25.3
Swap usage: 0%

Expanded Security Maintenance for Applications is not enabled.

0 updates can be applied immediately.

Enable ESM Apps to receive additional future security updates.
See https://ubuntu.com/esm or run: sudo pro status

The list of available updates is more than a week old.
To check for new updates run: sudo apt update

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

```

- ResourceManager: 52.207.250.246 (ec2-52-207-250-246.compute-1.amazonaws.com)

```

Problems Output Debug Console Terminal Ports
● cv_project scp -i ~/Downloads/ResourceManager.pem -r utils provider ubuntu@52.207.250.246:~/app/
config.py
init_.py
config_cpython-313.pyc
calculations_cpython-313.pyc
__init__.cpython-313.pyc
calculations.py
resource_manager.py
__init__.py
resource_manager_cpython-313.pyc
__init__.cpython-313.pyc
provider_app.py
● + cv_project

```

Time	File	Size	Speed	Progress
100%	3081	9.1KB/s	00:00	ssh
100%	51	0.2KB/s	00:00	ssh
100%	3106	9.6KB/s	00:00	ssh
100%	11KB	32.0KB/s	00:00	ssh
100%	204	0.7KB/s	00:00	ssh
100%	10KB	30.8KB/s	00:00	ssh
100%	7168	20.7KB/s	00:00	ssh
100%	48	0.1KB/s	00:00	ssh
100%	8700	25.2KB/s	00:00	ssh
100%	204	0.6KB/s	00:00	ssh
100%	6666	20.3KB/s	00:00	ssh

Screenshot of the AWS EC2 Instances page for instance i-0f963f5e6b3f8705e.

Instance summary for i-0f963f5e6b3f8705e (ResourceManager) Info

Updated less than a minute ago

Instance ID	i-0f963f5e6b3f8705e	Public IPv4 address	52.207.250.246 open address
IPv6 address	-	Instance state	Running
Hostname type	IP name: ip-172-31-24-148.ec2.internal	Private IP DNS name (IPv4 only)	ip-172-31-24-148.ec2.internal
Answer private resource DNS name	IPv4 (A)	Instance type	t2.micro
Auto-assigned IP Address	52.207.250.246 [Public IP]	VPC ID	vpc-0d423e8abb901bba7
IAM Role	LabRole	Subnet ID	subnet-034e45382800e5767
IMDSv2	Required	Instance ARN	arn:aws:ec2:us-east-1:210058569281:instance/i-0f963f5e6b3f8705e
Operator	-	AWS Compute Optimizer finding	Opt-in to AWS Compute Optimizer for recommendations. Learn more

Details **Status and alarms** **Monitoring** **Security** **Networking** **Storage** **Tags**

Instance details **Info**

AMI ID	ami-0ecb62995f68bb549	Monitoring	disabled	Platform details	Linux/UNIX
AMI name	ubuntu/images/hvm-ssd-gp3/ubuntu-noble-24.04-amd64-server-20251022	Allowed image	-	Termination protection	Disabled
Stop protection	Disabled	Launch time	Sat Jan 17 2026 18:39:25 GMT+0500 (Pakistan Standard Time) (1 min ago)	AMI location	amazon/ubuntu/images/hvm-ssd-gp3/ubuntu-noble-24.04-amd64-server-20251022

© 2026, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Terminal

```
229.9/229.9 kB 41.8 MB/s eta 0:00:00
Downloading six-1.17.0-py2.py3-none-any.whl (11 kB)
Installing collected packages: urllib3, six, numpy, markupsafe, jmespath, itsdangerous, idna, click, charset_normalizer, certifi, blinker, werkzeug, requests, python-dateutil, jinja2, Flask, botocore >=3transfer, boto3, S3Transfer, botocore >=3transfer, boto3 >=3transfer, botocore >=3.0.0 boto3 >=1.42.30 botocore >=1.42.30 certifi >=2026.1.4 charset_normalizer >=3.4.4 click >=8.3.1 idna >=3.11 itsdangerous >=2.2.0 jinja2 >=3.1.6 jmespath >=1.0.1 markupsafe >=3.0.3 numpy >=2.4.1 python-dateutil >=2.9.0.post0 requests >=3.2.5 s3transfer >=0.16.0 six >=1.17.0 urllib3 >=2.6.3 werkzeug >=3.1.5
(venv) ubuntu@ip-172-31-24-148:~$ mkdir p ~/app
(venv) ubuntu@ip-172-31-24-148:~$ cd ~/app
(venv) ubuntu@ip-172-31-24-148:~/app$ ls -la
total 16
drwxr-xr-x 4 ubuntu ubuntu 4096 Jan 17 13:46 .
drwxr-xr-x 6 ubuntu ubuntu 4096 Jan 17 13:45 ..
drwxr-xr-x 3 ubuntu ubuntu 4096 Jan 17 13:45 provider
drwxr-xr-x 3 ubuntu ubuntu 4096 Jan 17 13:46 tools
(venv) ubuntu@ip-172-31-24-148:~/app$ export PORT=5001
(venv) ubuntu@ip-172-31-24-148:~/app$ nohup python provider/provider_app.py > provider.log 2>&1 &
[1] 3181
(venv) ubuntu@ip-172-31-24-148:~/app$ cat provider.log
nohup: ignoring input
==== Resource Manager Initialized ====
Managing 3 users and 5 resources
=====
Starting Resource Manager Flask App
Port: 5001
=====

* Serving Flask app 'provider_app'
* Debug mode on
WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
* Running on all addresses (0.0.0.0)
* Running on http://127.0.0.1:5001
* Running on http://172.31.24.148:5001
Press CTRL+C to quit
* Restarting with stat
* Debugger is active!
* Debugger PIN: 118-597-904
(venv) ubuntu@ip-172-31-24-148:~/app$
```

cursor agent to run Agent - 30K to generate command

Overall:

Screenshot of the AWS EC2 Instances page showing a list of 6 instances.

Instances (6) Info

Find Instance by attribute or tag (case-sensitive)

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IPv4 address
mondial-api	i-0a32099feed34011d	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1b	ec2-35-175-190-149.com...	35.175.190.149
my-first-ec2-i...	i-0bee105d8fc7dad9	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1b	ec2-52-91-90-67.comp...	52.91.90.67
S1-User	i-0a5d3fa0e1669b48c	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1b	ec2-18-207-117-128.com...	18.207.117.128
S2-User	i-0e00e90103a3d52c0	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1b	ec2-50-19-166-35.com...	50.19.166.35
ResourceManag...	i-0f963f5e6b3f8705e	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1b	ec2-52-207-250-246.co...	52.207.250.246
S3-User	i-05e20216b65e8bc9	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1b	ec2-98-80-12-101.com...	98.80.12.101

DynamoDB Table: ResourceAllocations

- Table Name: ResourceAllocations
- Partition Key: user_id (String)
- Sort Key: timestamp (String)
- Region: us-east-1
- Table ARN: arn:aws:dynamodb:us-east-1:210058569281:table/ResourceAllocations

The screenshot shows the AWS DynamoDB console interface. At the top, there's a success message: "The ResourceAllocations table was created successfully." Below this, the "Tables" section lists three tables: Leaderboard, ResourceAllocations, and Users. The ResourceAllocations table is selected. On the right, the detailed view for the ResourceAllocations table is shown.

General Information:

- Partition key: user_id (String)
- Sort key: timestamp (String)
- Capacity mode: On-demand
- Table status: Active
- Table size: 0 bytes

Read/write capacity:

- Capacity mode: On-demand
- Maximum read request units: -
- Maximum write request units: -

Additional Info:

- ARN copied: arn:aws:dynamodb:us-east-1:210058569281:table/ResourceAllocations

SQS Queue: SubmissionsQueue

- Queue Name: SubmissionsQueue
- Type: Standard Queue
- Queue URL: <https://sqs.us-east-1.amazonaws.com/210058569281/SubmissionsQueue>
- Region: us-east-1
- Messages processed: Multiple (handled by Lambda 1 trigger)

The screenshot shows the AWS SQS console with the SubmissionsQueue details page. At the top, there is a success message: "Queue SubmissionsQueue created successfully. You can now send and receive messages." Below this, the queue name is listed as "SubmissionsQueue". The "Type" is shown as "Standard Queue" with a "URL copied" link. The ARN is "arn:aws:sqs:us-east-1:210058569281:SubmissionsQueue". The "Dead-letter queue" is set to "-". There are tabs for "Queue policies", "Monitoring", "SNS subscriptions", "Lambda triggers", "EventBridge Pipes", "Dead-letter queue", "Tagging", "Encryption", and "Dead-letter queue redrive tasks". Under "Access policy", the JSON policy is displayed:

```
{  
  "Version": "2012-10-17",  
  "Id": "..._default_policy_ID",  
  "Statement": [  
    {  
      "Sid": "..._owner_statement",  
      "Effect": "Allow",  
      "Principal": {  
        "ARN": "arn:aws:iam::210058569281:root"  
      },  
      "Action": "SQS:*",  
      "Resource": "arn:aws:sqs:us-east-1:210058569281:SubmissionsQueue"  
    }  
  ]  
}
```

Lambda Functions:

- **submission-handler** (Lambda 1):
 - Function Name: submission-handler
 - Runtime: Python 3.11
 - Trigger: SQS Queue (SubmissionsQueue)
 - Environment Variables: DYNAMODB_TABLE, LAMBDA_FUNCTION_2, AWS_REGION
 - Role: Uses LabRole with DynamoDB, SQS, and Lambda permissions

The screenshot shows the AWS Lambda console. The top navigation bar includes 'Lambda > Functions > submission-handler'. The main area displays the 'Function overview' for 'submission-handler'. It shows a green success message: 'The trigger SubmissionsQueue was successfully added to function submission-handler.' Below this, there's a 'Function ARN' section with the value 'arn:aws:lambda:us-east-1:210058569281:function:submission-handler'. The 'Configuration' tab is selected, showing a 'Triggers (1)' section with one entry: 'SQS:SubmissionsQueue' (state: Creating). The 'Code' tab is also visible at the bottom.

- **get-and-send (Lambda 2):**
 - Function Name: get-and-send
 - Runtime: Python 3.11
 - Trigger: Invoked by Lambda 1
 - Environment Variables: DYNAMODB_TABLE, RESOURCE_MANAGER_URL, AWS_REGION
 - Timeout: 30 seconds

The screenshot shows the AWS Lambda console. The top navigation bar includes 'Lambda > Functions > get-and-send'. A green success message 'Successfully updated the function get-and-send.' is displayed. The function details show it was last modified '1 minute ago'. The 'Code source' tab is selected, displaying the Python code for the 'get_and_send' function:

```

    #!/usr/bin/python3
    # Amazon Lambda Function 2: Get and Send
    # Retrieves allocations from DynamoDB and sends to Resource Manager
    # Environment variables
    # DYNAMODB_TABLE = os.environ.get('DYNAMODB_TABLE', 'ResourceAllocations')
    # RESOURCE_MANAGER_URL = os.environ.get('RESOURCE_MANAGER_URL', '')
    # REGION = os.environ.get('AWS_REGION', 'us-east-1')
    # Initialize AWS clients
    dynamodb = boto3.resource('dynamodb', region_name=REGION)
    table = dynamodb.Table(DYNAMODB_TABLE)
    USER_IDS = ['S1', 'S2', 'S3']

```

4.2 Key Code Implementation

Utility Calculation Function:

```
def calculate_utility(allocation_vector, execution_times, resource_prices,
                      weight_time, weight_expense):
    time_vector = calculate_execution_time_vector(allocation_vector,
                                                   execution_times)
    expense_vector = calculate_expense_vector(allocation_vector,
                                                execution_times, resource_prices)

    max_time = max(time_vector)
    total_expense = sum(expense_vector)

    denominator = weight_time * max_time + weight_expense * total_expense
    return 1.0 / denominator if denominator > 0 else 0
```

Multiplexing Calculation:

```
def calculate_actual_execution_matrix(allocation_matrix,
                                      base_execution_times):
    for i in range(n_tasks):
        for j in range(m_resources):
            multiplexing_factor = sum(allocation_matrix[k][j] for k in
                                       range(n_tasks))
            if allocation_matrix[i][j] == 1:
                actual_time = multiplexing_factor *
                base_execution_times[i][j]
```

5. Results

```

Problems Output Debug Console Terminal Ports
=====
CLoud Resource Allocation - COMPLETE TEST
=====

STEP 1: INDEPENDENT OPTIMIZATION
=====

### PHASE 1.1: User Optimization (Independent) ###

--- Optimizing S1 ---
--- Initialized Optimizer for S1 ---
Subtasks: 2
Weights: wt=0.88, we=0.12
Constraints: Ts=80s, Ms=90s
--- Optimizing S1 (Step 1) ---
Execution times: [5.0, 4.2, 3.6, 3.0, 2.8]
Evaluating 10 possible allocations...
Feasible allocations: 1/10
Optimal allocation: (0, 0, 0, 1, 1)
Expected utility: 0.2525
S1: (0, 0, 0, 1, 1) - utility = 0.2525

--- Optimizing S2 ---
--- Initialized Optimizer for S2 ---
Subtasks: 3
Weights: wt=0.89, we=0.11
Constraints: Ts=80s, Ms=90s
--- Optimizing S2 (Step 1) ---
Execution times: [6.0, 5.0, 4.0, 3.5, 3.0]
Evaluating 10 possible allocations...
Feasible allocations: 1/10
Optimal allocation: (0, 0, 1, 1, 1)
Expected utility: 0.1794
S2: (0, 0, 1, 1, 1) - utility = 0.1794

--- Optimizing S3 ---
--- Initialized Optimizer for S3 ---
Subtasks: 3
Weights: wt=0.90, we=0.10
Constraints: Ts=80s, Ms=90s
--- Optimizing S3 (Step 1) ---
Execution times: [4.0, 3.5, 3.2, 2.8, 2.4]
Evaluating 5 possible allocations...
Feasible allocations: 5/5
Optimal allocation: (0, 1, 1, 1, 1)
Expected utility: 0.1986
S3: (0, 1, 1, 1, 1) - utility = 0.1986

### PHASE 1.2: Provider Processing (With Multiplexing) ###

--- Resource Manager Initialized ---
Managing 3 users and 5 resources
--- Processing Allocations ---
Allocation Matrix:
S1: [0, 0, 0, 1, 1]
S2: [0, 0, 1, 1, 1]
S3: [0, 1, 1, 1, 1]

Problems Output Debug Console Terminal Ports
=====
Actual Execution Time Matrix (t[i]):
Task 1: [0.00, 0.00, 0.00, 9.00, 8.40]
Task 2: [0.00, 0.00, 8.00, 10.50, 9.00]
Task 3: [0.00, 3.50, 6.40, 8.40, 7.20]

Expense Matrix (e[i]):
Task 1: [0.00, 0.00, 0.00, 5.40, 5.60]
Task 2: [0.00, 0.00, 6.00, 6.30, 6.00]
Task 3: [0.00, 4.20, 4.00, 5.04, 4.80]

S1 Results:
Max time: 9.00s
Total expense: 11.00€

S2 Results:
Max time: 10.50s
Total expense: 18.30€

S3 Results:
Max time: 8.40s
Total expense: 18.84€

--- Multiplexed Resources ---
R3: 2 users (S2, S3)
R4: 3 users (S1, S2, S3)
R5: 3 users (S1, S2, S3)

### PHASE 1.3: Actual Utility Calculation ###

--- Actual Utility for S1 ---
Time vector: ['0.00', '0.00', '0.00', '9.00', '8.40']
Expense vector: ['0.00', '0.00', '0.00', '5.40', '5.60']
Actual utility: 0.1082
Utility loss: 0.1443

--- Actual Utility for S2 ---
Time vector: ['0.00', '0.00', '8.00', '10.50', '9.00']
Expense vector: ['0.00', '0.00', '6.00', '6.30', '6.00']
Actual utility: 0.0860
Utility loss: 0.0514

--- Actual Utility for S3 ---
Time vector: ['0.00', '3.50', '6.40', '9.40', '7.20']
Expense vector: ['0.00', '4.20', '4.00', '5.04', '4.80']
Actual utility: 0.1059
Utility loss: 0.0928

===== STEP 1 RESULTS =====

User Allocation Vector Expected U Actual U Loss
S1 (0, 0, 0, 1, 1) 0.2525 0.1082 0.1443
S2 (0, 0, 1, 1, 1) 0.1794 0.0860 0.0514
S3 (0, 1, 1, 1, 1) 0.1986 0.1059 0.0928

Total Expected Utility: 0.6396
Total Actual Utility: 0.3022
Total Utility Loss: 0.3285

===== STEP 2: STRATEGY UPDATE WITH UPDATED EXECUTION TIMES =====

```

```

### PHASE 2.1: Calculate Updated Execution Times ###

--- Updated Execution Times for Step 2 ---

New Base Execution Times (t_new):
Task 1: [5.00, 5.37, 6.40, 12.30, 11.00]
Task 2: [6.00, 6.17, 6.80, 12.80, 11.20]
Task 3: [4.00, 4.67, 8.00, 12.10, 10.60]

### PHASE 2.2: User Re-optimization ###

--- Re-optimizing S1 with updated times ---

--- Optimizing S1 (Step 1) ---
Execution times: [5.0, 5.3666666666666667, 6.4, 12.29999999999999, 11.0]
Evaluating 10 possible allocations...
Feasible allocations: 3/10
Optimal allocation: (1, 1, 0, 0, 0)
Expected utility: 0.1641
S1: (1, 1, 0, 0, 0) - utility = 0.1641

--- Re-optimizing S2 with updated times ---

--- Optimizing S2 (Step 1) ---
Execution times: [6.0, 6.1666666666666667, 6.8, 12.79999999999999, 11.2]
Evaluating 10 possible allocations...
Feasible allocations: 1/10
Optimal allocation: (1, 1, 1, 0, 0)
Expected utility: 0.0930
S2: (1, 1, 1, 0, 0) - utility = 0.0930

--- Re-optimizing S3 with updated times ---

--- Optimizing S3 (Step 1) ---
Execution times: [4.0, 4.6666666666666667, 8.0, 12.09999999999998, 10.6]
Evaluating 5 possible allocations...
WARNING: No feasible allocation found for S3!
Feasible allocations: 0/5
Optimal allocation: (1, 1, 1, 1, 0)
Expected utility: 0.0657
S3: (1, 1, 1, 1, 0) - utility = 0.0657

### PHASE 2.3: Provider Processing (Step 2) ###

--- Processing Allocations ---

Allocation Matrix:
S1: (1, 1, 0, 0, 0)
S2: (1, 1, 0, 0, 0)
S3: (1, 1, 1, 1, 0)

Actual Execution Time Matrix (tij):
Task 1: [15.00, 12.60, 0.00, 0.00, 0.00]
Task 2: [16.00, 15.00, 8.00, 0.00, 0.00]
Task 3: [12.00, 10.50, 6.40, 2.80, 0.00]

Expense Matrix (eij):
Task 1: [5.00, 5.04, 0.00, 0.00, 0.00]
Task 2: [6.00, 6.00, 6.00, 0.00, 0.00]
Task 3: [4.00, 4.20, 4.80, 5.04, 0.00]

S1 Results:
Max time: 15.00s
Total expense: 10.04€

```

```

Problems Output Debug Console Terminal Ports zsh + □ ⊞ ... 

--- Multiplexed Resources ---
R1: 3 users (S1, S2, S3)
R2: 3 users (S1, S2, S3)
R3: 2 users (S2, S3)

### PHASE 2.4: Step 2 Actual Utilities ###

--- Actual Utility for S1 ---
Time vector: ['15.00', '12.60', '0.00', '0.00', '0.00']
Expense vector: ['5.00', '5.04', '0.00', '0.00', '0.00']
Actual utility: 0.0944
Utility loss: 0.0946

--- Actual Utility for S2 ---
Time vector: ['18.00', '15.00', '8.00', '0.00', '0.00']
Expense vector: ['6.00', '6.00', '6.00', '0.00', '0.00']
Actual utility: 0.0793
Utility loss: 0.0374

--- Actual Utility for S3 ---
Time vector: ['12.00', '10.50', '6.40', '2.80', '0.00']
Expense vector: ['4.00', '4.20', '4.80', '5.04', '0.00']
Actual utility: 0.0793
Utility loss: -0.0137

===== STEP 2 RESULTS =====

User Allocation Vector Expected U Actual U Loss
S1 (1, 1, 0, 0, 0) 0.1641 0.0944 0.0946
S2 (1, 1, 1, 0, 0) 0.0930 0.0793 0.0374
S3 (1, 1, 1, 1, 0) 0.0657 0.0793 -0.0137

Total Expected Utility: 0.3227
Total Actual Utility: 0.2643
Total Utility Loss: 0.1184

===== COMPARISON: STEP 1 vs STEP 2 =====

User Step 1 Actual U Step 2 Actual U Improvement
S1 0.1082 0.0944 -0.0388
S2 0.0880 0.0793 -0.0325
S3 0.1059 0.0793 -0.0265

Total 0.3022 0.2043 -0.0978

===== CONSTRAINT VERIFICATION =====

Step 1 Constraints:
S1: Time √ (9.00 ≤ 500), Budget √ (11.00 ≤ 28)
S2: Time √ (10.50 ≤ 500), Budget √ (18.00 ≤ 30)
S3: Time √ (6.40 ≤ 500), Budget √ (18.84 ≤ 36)

Step 2 Constraints:
S1: Time √ (15.00 ≤ 500), Budget √ (10.04 ≤ 28)
S2: Time √ (18.00 ≤ 500), Budget √ (18.00 ≤ 30)
S3: Time √ (12.00 ≤ 500), Budget √ (18.84 ≤ 36)

```

5.1 Step 1 Results: Independent Optimization

Step 1 Allocation Table

User	Allocation Vector	Expected Utility	Actual Utility	Utility Loss
S1	(0, 0, 0, 1, 1)	0.2525	0.1082	0.1443
S2	(0, 0, 1, 1, 1)	0.1794	0.0880	0.0914
S3	(0, 1, 1, 1, 1)	0.1986	0.1059	0.0928
Total	-	0.6306	0.3022	0.3285

Step 1 Execution Time Matrix (with multiplexing)

Task	R1	R2	R3	R4	R5
S1	0.00	0.00	0.00	9.00	8.40
S2	0.00	0.00	8.00	10.50	9.00
S3	0.00	3.50	6.40	8.40	7.20

Step 1 Expense Matrix

Task	R1	R2	R3	R4	R5
S1	0.00	0.00	0.00	5.40	5.60
S2	0.00	0.00	6.00	6.30	6.00
S3	0.00	4.20	4.80	5.04	4.80

Step 1 Resource Multiplexing

Resources with multiple users:

- R3: 2 users (S2, S3)
- R4: 3 users (S1, S2, S3)

- R5: 3 users (S1, S2, S3)

5.2 Step 2 Results: Strategy Update

Updated Execution Times for Step 2

Formula: $t_{\text{new_ij}} = \hat{t}_{\text{ij}} + \sum(t_{\text{ij}})/n$

Task	R1	R2	R3	R4	R5
S1	5.00	5.37	8.40	12.30	11.00
S2	6.00	6.17	8.80	12.80	11.20
S3	4.00	4.67	8.00	12.10	10.60

Step 2 Allocation Table

User	Allocation Vector	Expected Utility	Actual Utility	Utility Loss
S1	(1, 1, 0, 0, 0)	0.1641	0.0694	0.0946
S2	(1, 1, 1, 0, 0)	0.0930	0.0556	0.0374
S3	(1, 1, 1, 1, 0)	0.0657	0.0793	-0.0137
Total	-	0.3227	0.2043	0.1184

Step 2 Execution Time Matrix (with multiplexing)

Task	R1	R2	R3	R4	R5
S1	15.00	12.60	0.00	0.00	0.00
S2	18.00	15.00	8.00	0.00	0.00
S3	12.00	10.50	6.40	2.80	0.00

Step 2 Expense Matrix

Task	R1	R2	R3	R4	R5
S1	5.00	5.04	0.00	0.00	0.00
S2	6.00	6.00	6.00	0.00	0.00
S3	4.00	4.20	4.80	5.04	0.00

Step 2 Resource Multiplexing

Resources with multiple users:

- R1: 3 users (S1, S2, S3)
- R2: 3 users (S1, S2, S3)
- R3: 2 users (S2, S3)

5.3 Comparison: Step 1 vs Step 2

User	Step 1 Actual Utility	Step 2 Actual Utility	Improvement
S1	0.1082	0.0694	-0.0388
S2	0.0880	0.0556	-0.0325
S3	0.1059	0.0793	-0.0265
Total	0.3022	0.2043	-0.0978

5.4 Constraint Verification

Step 1 Constraints

User	Max Time	Deadline	Status	Total Expense	Budget	Status
S1	9.00	500	✓	11.00	20	✓
S2	10.50	300	✓	18.30	30	✓
S3	8.40	800	✓	18.84	30	✓

Step 2 Constraints

User	Max Time	Deadline	Status	Total Expense	Budget	Status
S1	15.00	500	✓	10.04	20	✓
S2	18.00	300	✓	18.00	30	✓
S3	12.00	800	✓	18.04	30	✓

6. Analysis

6.1 Impact of Resource Multiplexing

In Step 1, users optimized independently without considering that other users might select the same resources. When the Resource Manager calculated actual execution times, multiplexing was discovered:

- **Multiplexed Resources:** R3 (2 users: S2, S3), R4 (3 users: S1, S2, S3), R5 (3 users: S1, S2, S3)
- **Impact:** Execution times increased proportionally to the number of users sharing each resource. For example, R4 had a base execution time of 3.0s for S1, but with 3 users sharing it, the actual execution time became 9.00s ($3 \times 3.0\text{s}$).
- **Utility Loss:** Users experienced significant utility loss with an average of 0.1095 per user (total loss: 0.3285). S1 had the largest utility loss (0.1443) due to multiplexing on R4 and R5, which were its chosen resources.

Specific Impact:

- S1 chose R4 and R5, both shared by all 3 users, resulting in execution times of 9.00s and 8.40s respectively ($3\times$ the base times).
- S2 and S3 both shared R3, R4, and R5, causing their execution times to be multiplied accordingly.
- Despite meeting all constraints, the utility loss demonstrates the importance of considering resource contention in optimization.

6.2 Effectiveness of Step 2 Strategy Update

The formula $t_{\text{new_ij}} = \hat{t}_{\text{ij}} + \sum(t_{ij})/n$ incorporates knowledge from Step 1 into users' optimization for Step 2.

Observations:

- **Utility Changes:** None of the users improved their actual utility in Step 2. All three users experienced decreases: S1 (-0.0388), S2 (-0.0325), and S3 (-0.0265). Total system utility decreased from 0.3022 to 0.2043 (-0.0978).
- **Strategy Changes:** All users completely changed their allocation strategies. S1 moved from (0,0,0,1,1) to (1,1,0,0,0), avoiding the heavily multiplexed R4 and R5 in favor of R1 and R2. S2 changed from (0,0,1,1,1) to (1,1,1,0,0), also avoiding R4 and R5. S3 changed from (0,1,1,1,1) to (1,1,1,1,0), avoiding R5.
- **Resource Avoidance:** Resources R4 and R5, which were heavily multiplexed in Step 1 (3 users each), were largely avoided in Step 2. Instead, users shifted to R1 and R2, which became the new bottlenecks with 3 users each.
- **Paradoxical Result:** Despite incorporating information from Step 1, the total utility decreased. This occurred because users' updated strategies created new contention points (R1 and R2 became heavily multiplexed), and the updated execution time formula penalized resources that were previously popular, potentially leading to suboptimal global outcomes.
- **S3 Exception:** S3's expected utility (0.0657) was lower than its actual utility (0.0793), indicating a negative utility loss. This suggests S3 benefited from the actual allocation pattern compared to its expectation, though this may be due to the infeasible allocation warning (no feasible allocation was found, so the optimizer selected the best infeasible option).

6.3 Game-Theoretic Insights

This two-step process demonstrates:

1. **Nash Equilibrium Seeking:** Users adjust strategies based on observed outcomes. In Step 2, all users avoided the resources that caused high multiplexing in Step 1 (R4 and R5), demonstrating strategic adaptation.
2. **Strategic Learning:** Step 2 incorporates information from Step 1 through the updated execution time formula. However, this led to all users converging on similar strategies (primarily R1 and R2), creating new bottlenecks.
3. **Fairness vs Efficiency:** The results show a trade-off between individual optimization and system-wide efficiency. While individual users attempted to optimize their utilities based on learned information, the total system utility actually decreased. This illustrates the challenge of reaching a Nash equilibrium where no user can improve by unilaterally changing strategy, but the overall system performance may not be optimal. The constraint satisfaction was maintained for all users in both steps, ensuring feasibility, but efficiency decreased in Step 2.

7. Challenges and Solutions

7.1 Implementation Challenges

1. **Challenge:** IAM Role Creation Restrictions
 - **Issue:** AWS lab account had restricted IAM permissions, preventing direct creation or modification of IAM roles.
 - **Solution:** Used existing `LabRole` provided by the AWS lab environment. The role already had necessary permissions for DynamoDB, SQS, and Lambda operations. Documented this limitation in the implementation.
2. **Challenge:** DynamoDB Decimal Type Compatibility
 - **Issue:** Lambda functions using `boto3` to store data in DynamoDB encountered errors because DynamoDB doesn't support Python's native float types - it requires Decimal types.
 - **Solution:** Modified `submission_handler.py` to convert float values to Decimal type using `Decimal(str(expected_utility))` before storing in DynamoDB.
3. **Challenge:** Lambda Dependencies (`requests` library)
 - **Issue:** Lambda 2 needed the `requests` library, but it's not included in Python's standard library. Packaging with `pip install` failed due to network restrictions in the deployment environment.
 - **Solution:** Replaced `requests` library with Python's built-in `urllib.request` module, which is available in the Lambda runtime environment without additional dependencies.
4. **Challenge:** Lambda Handler Configuration
 - **Issue:** Lambda functions initially failed with "Unable to import module 'lambda_function'" errors because the handler was set to the default `lambda_function.lambda_handler` instead of the actual module names.
 - **Solution:** Updated Lambda function handler configuration to `submission_handler.lambda_handler` for Lambda 1 and `get_and_send.lambda_handler` for Lambda 2.
5. **Challenge:** DynamoDB Composite Key Updates
 - **Issue:** Lambda 2 failed to update item status in DynamoDB because it was only providing the partition key (`user_id`) but not the sort key (`timestamp`), which is required for composite keys.
 - **Solution:** Modified `get_pending_allocations()` to return both allocations and timestamps, then updated `update_allocation_status()` to use both keys in the DynamoDB update operation.
6. **Challenge:** JSON Serialization of Decimal Types

- **Issue:** DynamoDB returns Decimal types, which cannot be directly serialized to JSON when sending to the Resource Manager.
- **Solution:** Converted Decimal values to integers when retrieving allocation vectors from DynamoDB using list comprehension: `[int(v) for v in allocation_vector]`.

7.2 Testing Challenges

- 1. Local vs AWS Testing:**
 - **Challenge:** Needed to test locally before deploying to AWS to catch errors early.
 - **Solution:** Developed comprehensive local test suite (`run_tests.py complete`) that simulates the complete flow without AWS dependencies. This allowed validation of optimization algorithms, matrix calculations, and utility computations before deployment.
- 2. Debugging Distributed System:**
 - **Challenge:** Errors in distributed systems are harder to debug because components run on different machines and logs are scattered.
 - **Solution:** Used AWS CloudWatch logs for Lambda functions, Flask application logs on EC2 instances, and health check endpoints on all services. Added extensive print statements and error handling to trace the flow through each component.
- 3. Network Connectivity Testing:**
 - **Challenge:** Verifying that all EC2 instances could communicate with each other and with AWS services.
 - **Solution:** Used security groups to allow necessary ports (5000, 5001), tested connectivity with curl commands, and verified health endpoints before running the complete flow.
- 4. Lambda Function Testing:**
 - **Challenge:** Lambda functions run in a managed environment, making local debugging difficult.
 - **Solution:** Tested Lambda code logic locally with mock events before deployment, then used CloudWatch logs extensively to debug runtime issues. Created minimal deployment packages to reduce complexity.

8. Conclusion

This project successfully implemented a game-theoretic resource allocation system demonstrating:

1. **Technical Implementation:** Successfully deployed a distributed system using EC2, Lambda, SQS, and DynamoDB
2. **Optimization Algorithm:** Implemented brute-force optimization with constraint checking
3. **Two-Step Game:** Demonstrated how users can improve strategies with learned information
4. **Resource Multiplexing:** Illustrated real-world challenges of shared resource allocation

Key Findings:

- **Step 1 Performance:** Total expected utility was 0.6306, but actual utility after multiplexing was 0.3022, showing a utility loss of 0.3285 (52% reduction) due to resource contention. All constraints were satisfied.
- **Step 2 Performance:** Despite incorporating learned information, total actual utility decreased to 0.2043 (32% decrease from Step 1). This demonstrates that individual strategic adjustments can lead to worse global outcomes, illustrating the complexity of game-theoretic resource allocation.
- **Constraint Satisfaction:** All constraints (deadlines and budgets) were successfully met in both steps for all three users, validating the feasibility of the allocations.
- **Resource Multiplexing Impact:** Resources R4 and R5 were heavily contested in Step 1 (3 users each), leading to 3x execution time multipliers. Users strategically avoided these in Step 2, but this created new bottlenecks on R1 and R2.

9. References

1. Wei, G., Vasilakos, A. V., Zheng, Y., & Xiong, N. (2009). A game-theoretic method of fair resource allocation for cloud computing services. *The Journal of Supercomputing*, 54(2), 252-269.
2. Course Materials: HY452 Project 2025, University of Crete
3. AWS Documentation: EC2, Lambda, DynamoDB, SQS

Appendix

A. Complete Code Listings

Attached zip file pdf code with this pdf as well.

B. AWS Configuration Details

IAM Roles:

- **LabRole:** Used by Lambda functions and EC2 instances
 - Note: Due to AWS lab account restrictions, custom IAM roles could not be created. The existing LabRole was used, which provided necessary permissions for DynamoDB, SQS, and Lambda operations.
 - Permissions include: DynamoDB read/write, SQS send/receive, Lambda invoke, EC2 basic operations

Resource URLs and Identifiers:

- **SQS Queue:**
 - Name: SubmissionsQueue
 - URL: <https://sqs.us-east-1.amazonaws.com/210058569281/SubmissionsQueue>
 - Region: us-east-1
- **DynamoDB Table:**
 - Name: ResourceAllocations
 - ARN: arn:aws:dynamodb:us-east-1:210058569281:table/ResourceAllocations
 - Region: us-east-1
 - Keys: Partition Key (user_id), Sort Key (timestamp)
- **EC2 Instances:**
 - User S1: <http://18.207.117.128:5000>
 - User S2: <http://50.19.166.35:5000>
 - User S3: <http://98.80.12.101:5000>

- Resource Manager: <http://52.207.250.246:5001>
- Lambda Functions:
 - submission-handler: Processes SQS messages, stores in DynamoDB
 - get-and-send: Retrieves from DynamoDB, sends to Resource Manager
 - Region: us-east-1

Environment Variables:

Lambda 1 (submission-handler):

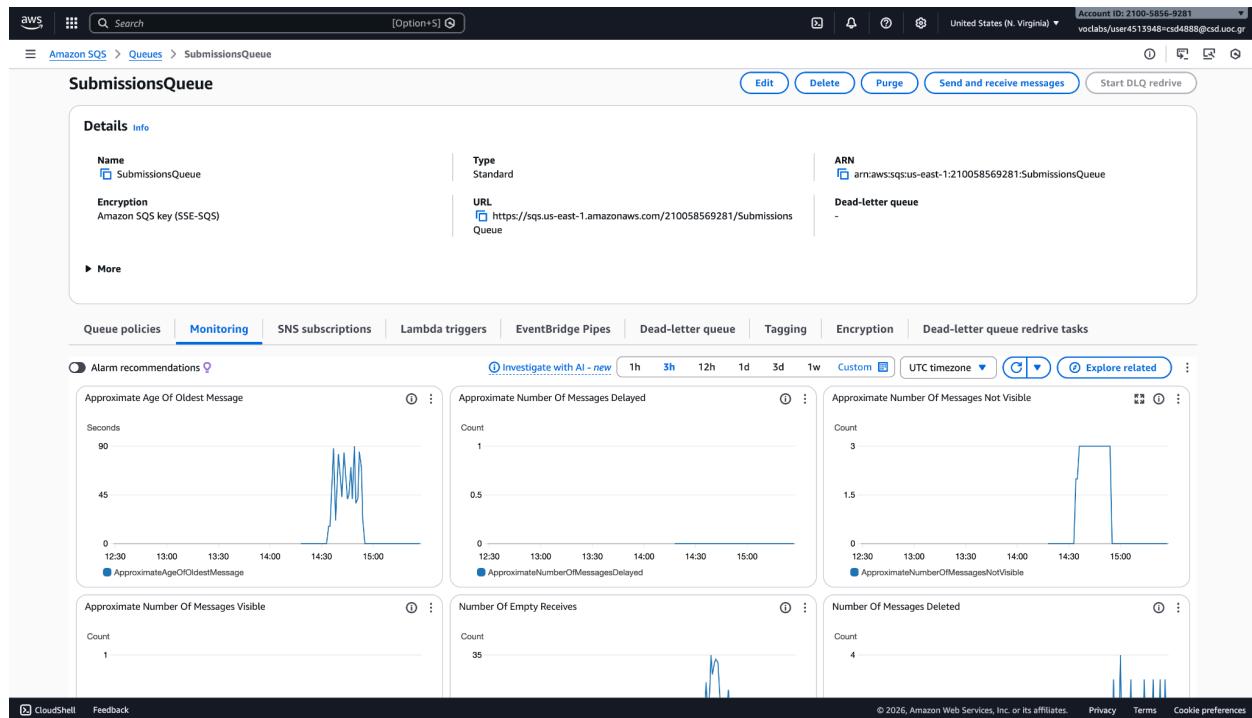
- DYNAMODB_TABLE: ResourceAllocations
- LAMBDA_FUNCTION_2: get-and-send
- AWS_REGION: us-east-1

Lambda 2 (get-and-send):

- DYNAMODB_TABLE: ResourceAllocations
- RESOURCE_MANAGER_URL: <http://52.207.250.246:5001>
- AWS_REGION: us-east-1 (automatically set by Lambda)

C. AWS Execution Logs

SQS QUEUE:



LAMBDA 1:

Diagram

Submission Handler

SQS

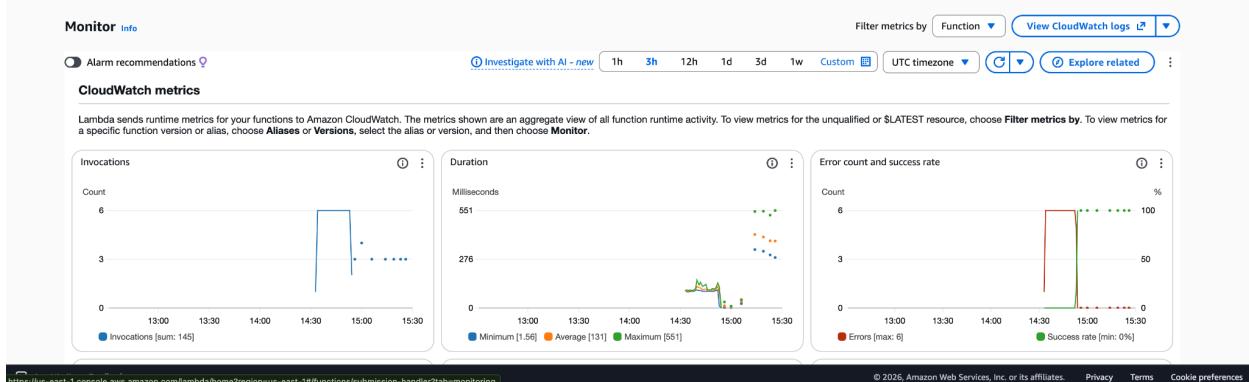
Monitor

CloudWatch metrics

Invocations

Duration

Error count and success rate



Log events

Timestamp Message

- 2026-01-17T15:26:20.809Z Step 1 - Submitted: {'\$2', '\$1', '\$3'}, All submitted: True
- 2026-01-17T15:26:20.809Z All users submitted for step 1, triggering Lambda 2
- 2026-01-17T15:26:21.071Z Triggered Lambda 2: 202
- 2026-01-17T15:26:21.073Z END RequestId: 56bf4242-2ac7-5154-b875-23a7868eb54e
- 2026-01-17T15:26:21.073Z REPORT RequestId: 56bf4242-2ac7-5154-b875-23a7868eb54e Duration: 551.09 ms Billed Duration: 552 ms Memory Size: 128 MB Max Memory Used: 89 MB
- 2026-01-17T15:26:21.289Z START RequestId: eed5b5d6-49b5-57e4-9e08-d443e5e5f658 Version: \$LATEST
- 2026-01-17T15:26:21.290Z Received event: {"Records": [{"messageId": "9fa5b1ef-za21-462e-a88b-b90eee946a7b", "receiptHandle": "AQEBF7HP6kbYcv109xMsx1Arsrz22CrzLdSoER/ksRvu83Nrxg0Qfqr..."}]
- 2026-01-17T15:26:21.290Z Processing allocation for S2, step 1
- 2026-01-17T15:26:21.389Z Stored allocation for S2 in DyanomoDB
- 2026-01-17T15:26:21.409Z Step 1 - Submitted: {'\$2', '\$1', '\$3'}, All submitted: True
- 2026-01-17T15:26:21.409Z All users submitted for step 1, triggering Lambda 2
- 2026-01-17T15:26:21.580Z Triggered Lambda 2: 202
- 2026-01-17T15:26:21.589Z END RequestId: eed5b5d6-49b5-57e4-9e08-d443e5e5f658
- 2026-01-17T15:26:21.589Z REPORT RequestId: eed5b5d6-49b5-57e4-9e08-d443e5e5f658 Duration: 299.44 ms Billed Duration: 300 ms Memory Size: 128 MB Max Memory Used: 89 MB
- 2026-01-17T15:26:22.103Z START RequestId: 211d8de7-5c15-5c1a-b2ae-ec02f04fc466 Version: \$LATEST
- 2026-01-17T15:26:22.103Z Received event: {"Records": [{"messageId": "273902cc-3b88-428a-8262-e9e48984d75a", "receiptHandle": "AQEBsb/fr004Hw3F1v0HDNG3AyEsd130Xicdf0W4P/pW+6kx0mNYowsZ..."}]
- 2026-01-17T15:26:22.108Z Processing allocation for S3, step 1
- 2026-01-17T15:26:22.209Z Stored allocation for S3 in DyanomoDB
- 2026-01-17T15:26:22.209Z Step 1 - Submitted: {'\$2', '\$1', '\$3'}, All submitted: True
- 2026-01-17T15:26:22.209Z All users submitted for step 1, triggering Lambda 2
- 2026-01-17T15:26:22.377Z Triggered Lambda 2: 202
- 2026-01-17T15:26:22.389Z END RequestId: 211d8de7-5c15-5c1a-b2ae-ec02f04fc466
- 2026-01-17T15:26:22.389Z REPORT RequestId: 211d8de7-5c15-5c1a-b2ae-ec02f04fc466 Duration: 285.56 ms Billed Duration: 286 ms Memory Size: 128 MB Max Memory Used: 89 MB

LAMBDA 2:

Last modified
8 minutes ago

Function ARN
arn:aws:lambda:us-east-1:210058569281:function:get-and-send

Function URL | Info

Code | Test | **Monitor** | Configuration | Aliases | Versions

CloudWatch metrics

Invocations

Count

Duration

Milliseconds

Error count and success rate

Count %

© 2026, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

CloudWatch

Favorites and recents

Ingestion

CloudWatch Logs

Log Management New

Log Anomalies

Live Tail

Logs Insights

Contributor Insights

Metrics

Network Monitoring

Setup

Log events

Filter events - press enter to search

Actions | Start tailing | Create metric filter

Timestamp | Message

No older events at this moment. *Retry*

2026-01-17T15:26:21.312Z INIT_START Runtime Version: python:3.11.v109 Runtime Version ARN: arn:aws:lambda:us-east-1:runtime:49f733259c7ce7e0dee75ff91c6afe35c7d58b04ed300f32701216263...

2026-01-17T15:26:21.790Z START RequestId: 3996d2ad-6f25-4c2d-b8f5-e7b57d5afca8 Version: \$LATEST

Received event: {"step": 1, "trigger_source": "submission_handler"}
Found allocation for S1: [0, 0, 0, 1, 1]
Found allocation for S2: [0, 0, 1, 1, 1]
Found allocation for S3: [0, 1, 2, 1, 1]
Retrieved allocations for step 1: ['S1', 'S2', 'S3']
Sending to Resource Manager: http://S2.207.250.246:5001/calculate_matrices
Payload: {"allocations": {"S1": [0, 0, 0, 1, 1], "S2": [0, 0, 1, 1, 1], "S3": [0, 1, 1, 1, 1]}, "step": 1}

2026-01-17T15:26:22.119Z REPORT RequestId: 3996d2ad-6f25-4c2d-b8f5-e7b57d5afca8 Duration: 489.23 ms Billed Duration: 963 ms Memory Size: 128 MB Max Memory Used: 86 MB Init Duration: 4...

2026-01-17T15:26:22.152Z Resource Manager response: 200

2026-01-17T15:26:22.204Z Updated status for S1 to processed

2026-01-17T15:26:22.239Z Updated status for S2 to processed

2026-01-17T15:26:22.259Z Updated status for S3 to processed

2026-01-17T15:26:22.279Z END RequestId: 3996d2ad-6f25-4c2d-b8f5-e7b57d5afca8

2026-01-17T15:26:22.428Z REPORT RequestId: 3996d2ad-6f25-4c2d-b8f5-e7b57d5afca8 Duration: 489.23 ms Billed Duration: 963 ms Memory Size: 128 MB Max Memory Used: 86 MB Init Duration: 4...

2026-01-17T15:26:22.429Z START RequestId: 63366e7e-66ce-4d56-9162-08871426ff93 Version: \$LATEST

Received event: {"step": 1, "trigger_source": "submission_handler"}
WARNING: No pending allocation found for S1 in step 1
WARNING: No pending allocation found for S2 in step 1
Found allocation for S3: [0, 1, 1, 1, 1]

2026-01-17T15:26:22.479Z END RequestId: 63366e7e-66ce-4d56-9162-08871426ff93 Duration: 91.44 ms Billed Duration: 92 ms Memory Size: 128 MB Max Memory Used: 86 MB

2026-01-17T15:26:22.499Z REPORT RequestId: 63366e7e-66ce-4d56-9162-08871426ff93 Duration: 91.44 ms Billed Duration: 92 ms Memory Size: 128 MB Max Memory Used: 86 MB

2026-01-17T15:26:22.528Z REPORT RequestId: 63366e7e-66ce-4d56-9162-08871426ff93 Duration: 91.44 ms Billed Duration: 92 ms Memory Size: 128 MB Max Memory Used: 86 MB

2026-01-17T15:26:22.520Z REPORT RequestId: 63366e7e-66ce-4d56-9162-08871426ff93 Duration: 91.44 ms Billed Duration: 92 ms Memory Size: 128 MB Max Memory Used: 86 MB

© 2026, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

DYNAMODB:

The screenshot shows the AWS DynamoDB console interface. On the left, there's a navigation sidebar with sections for DynamoDB (Dashboard, Tables, Explore items, PartiQL editor, Backups, Exports to S3, Imports from S3, Integrations, Reserved capacity, Settings) and DAX (Clusters, Subnet groups, Parameter groups, Events). The main area is titled "ResourceAllocations". It has a "Scan or query items" section with a "Scan" button selected. Below it is a table titled "Table: ResourceAllocations - Items returned (12)". The table has columns: user_id, timestamp (String), allocation_vector, expected_utility, status, and step. The data in the table is as follows:

	user_id...	timestamp (String)	allocation_vector	expected_utility	status	step
<input type="checkbox"/>	S2	2026-01-17T15:14:18...	[{"N": "0"}, {"N": "0...]	0.1794365691727974	pending	1
<input type="checkbox"/>	S2	2026-01-17T15:19:17...	[{"N": "0"}, {"N": "0...]	0.1794365691727974	pending	1
<input type="checkbox"/>	S2	2026-01-17T15:23:46...	[{"N": "0"}, {"N": "0...]	0.1794365691727974	pending	1
<input type="checkbox"/>	S2	2026-01-17T15:26:21...	[{"N": "0"}, {"N": "0...]	0.1794365691727974	processed	1
<input type="checkbox"/>	S1	2026-01-17T15:14:17...	[{"N": "0"}, {"N": "0...]	0.25252525252525...	pending	1
<input type="checkbox"/>	S1	2026-01-17T15:19:17...	[{"N": "0"}, {"N": "0...]	0.25252525252525...	pending	1
<input type="checkbox"/>	S1	2026-01-17T15:23:45...	[{"N": "0"}, {"N": "0...]	0.25252525252525...	pending	1